
Criticality in neural networks paper title

Paul Rozdeba

Department of Physics
University of California, San Diego
La Jolla, CA 92093
prozdeba@physics.ucsd.edu

Forrest Sheldon

Department of Physics
University of California, San Diego
La Jolla, CA 92093
fsheldon@physics.ucsd.edu

Abstract

It has been posited that biological neural networks, such as a brain, may naturally exist in critical states. We propose two mechanisms for signal transduction in two such networks as encoding strategies which are optimized by criticality. First, we examine compressive sensing in a 2-dimensional Ising model at or near its critical temperature. Secondly, we examine the dynamical synchronization capabilities of a random neural network model as it transitions into chaotic behavior. We propose that both techniques should be most successful at the critical state of either model.

1 Introduction

Criticality has enjoyed widespread success in the scientific community, being applied to problems as diverse as gene expression, starling flocks, and forest fires.[1] Originating in statistical physics, criticality is a set of properties of a system near a second order phase transition, in which the system possesses long range correlations in both space and time. These properties make it an attractive tool for describing the many complex systems that seem to display similar long range order. However, their analytical intractability restricts most arguments for criticality's existence to a rather superficial resemblance, usually resting on the existence of a power law in some measurement of the system. The recent realization that power laws occur far more often, and for a broader range of mechanisms than originally thought has tempered some of the enthusiasm for scale-free networks and, by association, criticality.[2] However, as the most established and flexible phenomenon displaying long range order, criticality still occupies a strong position as a candidate theory for understanding complex systems.

In the review "Emergent complex neural dynamics[1]" Dante Chialvo makes a case for the brain exhibiting criticality. He pulls from several pieces of evidence including:

- The brain contains the *necessary* elements (a large network of nonlinear interacting elements) to display complex emergent behavior and thus criticality.
- EEG/MEG readings of healthy brains do not show a preferred time scale.
- All models that display emergent complex behavior also display criticality.
- Neuronal avalanches have a scale free size distribution matching a critical branching process.
- Degree distributions created from fMRI recordings are scale-free.

Taken as a whole these form a reasonably strong case for criticality playing some role in the brain's dynamics however they all suffer from the defect mentioned previously: Power laws are nonspecific to criticality. As such, we take up the same investigation from the opposite direction asking 'Why would a brain want to be critical?' To this end, we examine systems that possess a known critical phase transition and which are relevant to neuroscience: the Ising model and a simple neural network with a leak current and sigmoidal connection strengths.

In the Ising model, we examine the structure of clustering that occurs in the vicinity of a phase transition. In particular, is it possible to reconstruct the states of distant spins held fixed as the system is evolved, given the state of the system at all points between them? We examine this possibility at various temperatures and examine the independence of multiple measurements.

2 The Ising Model

The Ising model could be considered the canonical physical system exhibiting critical phase transition. The spins may be considered small arrows that may point up or down with a corresponding value of +1 and -1. They are governed by the Hamiltonian,

$$H = -\frac{\epsilon}{2} \sum_{i \neq j} s_i s_j$$

where ϵ is some energy scale and the factor of $\frac{1}{2}$ compensates for double counting in the sum. There is an energetic cost for two neighboring spins to oppose each other. As such, the system displays two phases: As the temperature is raised, the system moves from an ordered phase in which all of the spins point in the same direction and thus minimize the energy, to a disordered phase in which neighboring spins oppose each other. Between the two phases the system displays power law correlation distributions in space and time and a large number of accessible states. These three regimes are displayed in Figure 1.




Figure 1: The Ising model below, near, and above the critical temperature

3 Compressive Sensing

Compressive sensing is a process by which a k -sparse signal of length N may be perfectly reconstructed from only $O(K \log N/K)$ measurements if those measurements are random projections.[3] This is interesting both because $O(K \log N/K)$ is much better than $O(N)$ which is implied by the Shannon-Nyquist theorem, and that this is done using random projections. This is because random projections are information preserving in that they approximately preserve distances between major features in the data. Stating this more mathematically, given a discrete signal x with a sparse representation in some basis Ψ , we acquire x by projecting it onto a set of vectors Φ , acquiring a vector of projections y :

$$y = \Phi x = \Phi \Psi s = \Theta s$$

We can perfectly reconstruct our original signal s from only a few measurements y so long as $\Theta = \Phi \Psi$ obeys the *Restricted Isometry Property* which happens to be fulfilled by most random bases. (Interested readers may consult[3][4]) The sparsity of most natural signals and the use of random projections has led Chuck Stevens to suggest that this coding strategy may be used in a modified form in several regions of the cortex that have resisted attempts at making a 'map' of their function. We propose that a system at criticality may naturally perform these projections and choose an Ising model to attempt to demonstrate it.

4 Methods-I

Ising models were simulated by means of Markov chain Monte Carlo. [5][6] Comparison between Metropolis and Gibbs sampling algorithms yielded that the Metropolis algorithm performed poorly

at high temperatures and Gibbs sampling was performed thereafter. A signal vector \vec{s} was generated with $[-1, 1]$ entries. These entries were inputted into the model at random points, held fixed as the system evolved. A single random spin r was also selected to be recorded. We would regard this single spin as our random projection of the signal held fixed in the model. After evolving the system to equilibrium, a breadth first search was performed through the cluster in the model containing the recording spin. A cluster incidence vector \vec{i} was formed whose elements were 1 if the corresponding signal element was a member of the cluster and zero otherwise. The signal vector was then l^1 normalized so that the equation, $r = \vec{i} \cdot \vec{s}$ was satisfied. By repeating this procedure with the signal spins held fixed and the same recording spin, we were able to generate a system $\vec{r} = I\vec{s}$ that should be approximately solved by our signal vector and recorded spins.

A full study of the properties of these measurement vectors (and alternative schemes at constructing a measurement matrix) is still underway. Under the limitation of time, two analyses were undertaken. First, the rank of I was examined as a function of number of measurements and temperature. Second, the system was solved by means of the Moore-Penrose pseudoinverse and the reconstruction error calculated, also as a function of number of measurements and temperature. All results here are for a (32×32) spin model to keep computation time reasonable, although similar analyses have been run on larger systems.

5 Results-I

Results thus far have been...a bit mundane. We begin with the high point. The rank of the measurement matrix (Figure 2) seems to show a distinct temperature dependence. At the critical temperature, each measurement remains independent and the measurement matrix attains full row rank in as few steps as possible. Only slightly worse is the low temperature limit. Here several measurement vectors are linearly dependent on those previously measured and do not increase the rank. Finally, worst off is the high temperature limit in which the matrix was not able to reach full rank given twice as many measurements as necessary. This would seem to indicate that the critical state offers some ad-

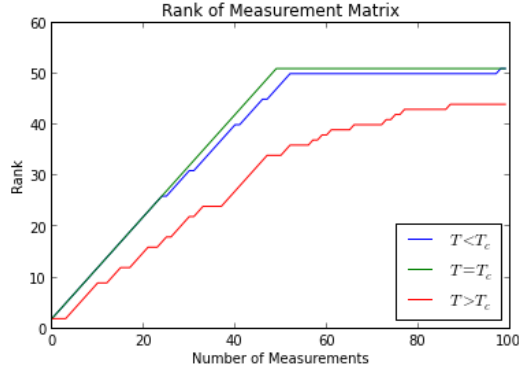


Figure 2: Measurement matrix ranks at various temperatures

vantage over its neighboring phases in its ability to transmit information located around the model. In the disordered phase, the small correlation length makes it very difficult to obtain interactions with distant spins, and in the ordered phase the large interaction size causes repetitive couplings to all the spins also transmitting redundant information.

Attempting to reconstruct the original signal from these measurements however, has been far less successful. Using the Moore-Penrose pseudoinverse to attempt reconstruction at every step, we were able to calculate reconstruction error for each measurement (Figure 3). We have far less understanding of this figure. It seems that the critical temperature is the most numerically unstable of all and especially so just as the matrix reaches full rank. As this is the case for the low temperature matrix as well, this may be a feature of the MP pseudoinverse. While an l^1 minimization would have been more enkeeping with the original inspiration for this project due to its preferred role in compressive sensing, it is substantially more difficult to implement and refinements of the reconstruction will

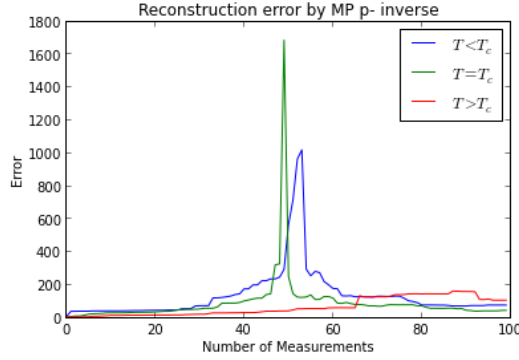


Figure 3: Reconstruction error at various temperatures

be relegated to subsequent work. As it stands, I believe the poor quality of the reconstructions is feature of the specific encoding and reconstruction scheme used and that further work may yield substantially more accurate results.

6 Dynamical synchronization technique

The model equations are coupled to a set of data through a linear term as

$$\frac{dx_i}{dt} = f_i(x) + \sum_j g_{ij} (y_j - x_j) \quad (1)$$

where $x(t)$ and $y(t)$ are the model and the data trajectories, respectively, and g_{ij} is a set of coupling constants. In practice, only a few of the couplings will be nonzero (corresponding to the components of the system which can be measured) and the matrix g_{ij} is diagonal, since in this scheme there is no obvious advantage to coupling different components of the system to each other.

In general, a model can be synchronized to a data set if enough measurements are made, and if the coupling is large enough. What constitutes “enough” seems to come on a case-by-case basis. Roughly, what is required is for the linear terms to regularize the (coupled) solution manifold by making the value of the largest Lyapunov exponent non-positive. However, the ratio of the i -th coupling term to the i -th term of the dynamics must remain small, so as to not wash out the physicality of the problem (in practice, this may also increase the speed of convergence to the actual solution). In other words, adjusting the coupling is a delicate balance for which there is no general procedure.

The success of the procedure, however, is not ultimately measured by the ability to synchronize, but by the ability of the model *make predictions* after the coupling is turned off, since in general we’re only able to measure a few components of the system. In the case where we have access to all of the data states, however, we can additionally test the procedure by examining the normally unobservable states. This is called a *twin experiment* in the literature and is a useful technique for assessing the viability of the model to synchronize in a controlled setting.

7 Random neural network model

We now consider a neural network model, originally presented by [7], which is a dynamical system describing a network of N randomly connected elements each having a leaky capacitive term. The equations of motion describing the system are

$$\frac{dV_i}{dt} = -V_i + \sum_{j=1}^N J_{ij} \phi(V_j) \quad (2)$$

where $\phi(V_i)$ is a function which may be thought of as an “activity” proportional to the synaptic current between neurons i and j . One possible choice is a sigmoid function of V , i.e. $\phi_i = \tanh(\alpha V_i)$.

This choice is both biologically motivated as acting to saturate synaptic activity as a function of membrane voltage, as well as mathematically to avoid highly unstable, runaway solutions to eq. (2). α acts as a control parameter on the turnaround rate of the synaptic activity around $V = 0$; in some sense it controls the “degree of nonlinearity” in the system.

The matrix J_{ij} describes the connectivity in the network; in this particular model, J_{ij} is chosen to be a Gaussian random matrix with elements distributed according to the statistics

$$\langle J_{ij} \rangle = 0, \quad \langle J_{ij} J_{kl} \rangle = \frac{\tilde{J}^2}{N} \delta_{ik} \delta_{jl} (1 - \delta_{ij}) (1 - \delta_{kl}) \quad (3)$$

which means synaptic connections are, in general, asymmetric and totally decorrelated. This also means that, on average, half of the connections are inherently inhibitory and half are excitatory. We use “inherently” because the sign of the synaptic current into neuron i switches depending on the sign of V_j .

A detailed mathematical treatment of this model in the limit of large N is given in [7]. Under the replacement $\alpha \rightarrow g\tilde{J}$ in the expression for $\phi(V_i)$, the model undergoes a “phase transition” when the control parameter $g\tilde{J}$ reaches a critical value of 1. This is manifest in the structure of the attracting solution manifold in the $\{V_i\}$ state space, which acquires a positive Lyapunov exponent; in other words a family of chaotic solutions to (2) appears.

The similarity of this transition to the phase transition in the 2D Ising model near T_{crit} is most apparent in the behavior of the quantity

$$\Delta(t) \equiv \langle V_i(t_0) V_i(t_0 + t) \rangle \quad (4)$$

7.1 Numerical analysis

We examined the model for $N = 256$ neurons with a single instantiation of J_{ij} . Ideally, we would like to scale up the simulation size to a larger N , perhaps near 1000, and to gather statistics about an *ensemble* of models parameterized by different instantiations of J_{ij} . However, limited by time, we now present said results as at least a preliminary examination of the model.

First, we performed a comparison of numerical results to the analytic results of [7]. For a single randomly drawn instantiation of J_{ij} , (2) was integrated forward in time using the LSODA solver in `scipy.integrate`. The initial conditions were drawn randomly from a ball of radius 1 near the origin $V_i = 0$. We found that above $g\tilde{J} = 1$, limit cycle solutions appeared in a Hopf bifurcation from an attracting fixed point at the origin (which became unstable in the bifurcation).

References

- [1] Dante R. Chialvo. Emergent complex neural dynamics. *Nature Physics*, 6(10):744–750, October 2010.
- [2] Evelyn Fox Keller. Revisiting “scale-free” networks. *BioEssays : news and reviews in molecular, cellular and developmental biology*, 27(10):1060–8, October 2005.
- [3] EJ Candès and MB Wakin. An introduction to compressive sampling. *Signal Processing Magazine, IEEE*, (March 2008):21–30, 2008.
- [4] RG Baraniuk. Compressive sensing. *Signal Processing Magazine, IEEE*, 24(July):1–9, 2007.
- [5] David J C MacKay. *Information Theory, Inference, and Learning Algorithms* David J.C. MacKay, volume 100. 2003.
- [6] Daniel V. Schroeder. *An Introduction to Thermal Physics*, 1999.
- [7] H Sompolinsky, A Crisanti, and H. Sommers. Chaos in random neural networks. *Physical Review Letters*, 61(3):259–262, July 1988.